

Homework 2

Yuki Joyama

2024-02-16

1

```
# load data
dose = c(0, 1, 2, 3, 4)
num = c(30, 30, 30, 30, 30)
dead = c(2, 8, 15, 23, 27)
data_1 = data.frame(dose, num, dead)

# visualization
# plot(data_1$dose, data_1$dead/data_1$num, xlab = 'dose', ylab = 'Proportion dying', cex = 1.5, pch = .)

# data prep
x = data_1$dose
y = data_1$dead
m = data_1$num
resp = cbind(y, m-y)
```

Now, I will fit the model $g(P(\text{dying})) = \alpha + \beta X$ with logit, probit, and complementary log-log links.

```
# fit logistic model, logit
glm_logit = glm(resp ~ x, family = binomial(link = 'logit'))
summary(glm_logit)

##
## Call:
## glm(formula = resp ~ x, family = binomial(link = "logit"))
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.3238      0.4179  -5.561 2.69e-08 ***
## x              1.1619      0.1814   6.405 1.51e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 64.76327  on 4  degrees of freedom
## Residual deviance:  0.37875  on 3  degrees of freedom
## AIC: 20.854
##
## Number of Fisher Scoring iterations: 4
```

```
# fit logistic model, probit
glm_probit = glm(resp ~ x, family = binomial(link = 'probit'))
summary(glm_probit)

##
## Call:
## glm(formula = resp ~ x, family = binomial(link = "probit"))
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.37709    0.22781  -6.045 1.49e-09 ***
## x           0.68638    0.09677   7.093 1.31e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 64.76327  on 4  degrees of freedom
## Residual deviance:  0.31367  on 3  degrees of freedom
## AIC: 20.789
##
## Number of Fisher Scoring iterations: 4
```

```
# fit logistic model, cloglog
glm_cloglog = glm(resp ~ x, family = binomial(link = 'cloglog'))
summary(glm_cloglog)

##
## Call:
## glm(formula = resp ~ x, family = binomial(link = "cloglog"))
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.9942    0.3126  -6.378 1.79e-10 ***
## x           0.7468    0.1094   6.824 8.86e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 64.7633  on 4  degrees of freedom
## Residual deviance:  2.2305  on 3  degrees of freedom
## AIC: 22.706
##
## Number of Fisher Scoring iterations: 5
```

(a)

```
# 95% CI for beta, logit
beta = glm_logit$coefficients[2]
```

```
se = sqrt(vcov(glm_logit)[2, 2])
round(beta + c(qnorm(0.025), -qnorm(0.025)) * se, 3)
```

```
## [1] 0.806 1.517
```

```
# 95% CI for beta, probit
beta = glm_probit$coefficients[2]
se = sqrt(vcov(glm_probit)[2, 2])
round(beta + c(qnorm(0.025), -qnorm(0.025)) * se, 3)
```

```
## [1] 0.497 0.876
```

```
# 95% CI for beta, cloglog
beta = glm_cloglog$coefficients[2]
se = sqrt(vcov(glm_cloglog)[2, 2])
round(beta + c(qnorm(0.025), -qnorm(0.025)) * se, 3)
```

```
## [1] 0.532 0.961
```

```
# p_hat(dying/x = 0.01), logit
predict(glm_logit, newdata = data.frame(x = 0.01), type = 'response')
```

```
##          1
## 0.09011997
```

```
# calculate by hand
# or = exp(coef(glm_logit)[1] + 0.01 * coef(glm_logit)[2])
# or / (1 + or)

# p_hat(dying/x = 0.01), probit
predict(glm_probit, newdata = data.frame(x = 0.01), type = 'response')
```

```
##          1
## 0.0853078
```

```
# p_hat(dying/x = 0.01), cloglog
predict(glm_cloglog, newdata = data.frame(x = 0.01), type = 'response')
```

```
##          1
## 0.1281601
```

Model	Estimate of beta	CI for beta	Deviance	p_hat(dying x=0.01)
logit	1.162	(0.806-1.517)	0.379	0.0901
probit	0.686	(0.497-0.876)	0.314	0.0853
c-log-log	0.747	(0.532-0.961)	2.231	0.128

The estimate of beta in the logit model represents the change in the log odds of the response variable for a one-unit change in the predictor variable (dose). The 95% CI for the estimate of beta provides a range

within which we can be 95% confident that the true value of beta lies. The deviance can be used to check the goodness of fit of the models and 0.314 in probit model indicates a better fit. $\hat{p}(\text{dying}|x = 0.01)$ gives a probability estimate given that the predictor variable x takes the value of 0.01. In logit and probit model, the probabilities are similar. However, in the c-log-log model, which employs an asymmetric link function, the estimated probabilities appears to be larger than the other two models.

(b)

Three models can be expressed as below:

$$\log\left(\frac{p}{1-p}\right) = \alpha + \beta X$$

$$\Phi^{-1}(p) = \alpha + \beta X$$

$$\log(-\log(1-p)) = \alpha + \beta X$$

We want to estimate x_0 such that $\alpha + \beta X = g(p = 0.5)$

Given $p=0.5$,

$$\text{Logit: } 0 = \alpha + \beta x_0 \rightarrow x_0 = -\frac{\alpha}{\beta}$$

$$\text{var}(\hat{x}_0) = \left(\frac{\partial x_0}{\partial \alpha}\right)^2 \text{var}(\hat{\alpha}) + \left(\frac{\partial x_0}{\partial \beta}\right)^2 \text{var}(\hat{\beta}) + 2\left(\frac{\partial x_0}{\partial \alpha}\right)\left(\frac{\partial x_0}{\partial \beta}\right) \text{cov}(\hat{\alpha}, \hat{\beta})$$

$$\rightarrow \text{var}(\hat{x}_0) = \frac{1}{\beta^2} \text{var}(\hat{\alpha}) + \frac{\alpha^2}{\beta^4} \text{var}(\hat{\beta}) - 2\frac{\alpha}{\beta^3} \text{cov}(\hat{\alpha}, \hat{\beta})$$

Probit:

$$0 = \alpha + \beta x_0 \rightarrow x_0 = -\frac{\alpha}{\beta}$$

$$\text{var}(\hat{x}_0) = \frac{1}{\beta^2} \text{var}(\hat{\alpha}) + \frac{\alpha^2}{\beta^4} \text{var}(\hat{\beta}) - 2\frac{\alpha}{\beta^3} \text{cov}(\hat{\alpha}, \hat{\beta})$$

C-loglog:

$$-0.367 = \alpha + \beta x_0 \rightarrow x_0 = -\frac{0.367 + \alpha}{\beta}$$

$$\text{var}(\hat{x}_0) = \left(\frac{\partial x_0}{\partial \alpha}\right)^2 \text{var}(\hat{\alpha}) + \left(\frac{\partial x_0}{\partial \beta}\right)^2 \text{var}(\hat{\beta}) + 2\left(\frac{\partial x_0}{\partial \alpha}\right)\left(\frac{\partial x_0}{\partial \beta}\right) \text{cov}(\hat{\alpha}, \hat{\beta})$$

$$\rightarrow \text{var}(\hat{x}_0) = \frac{1}{\beta^2} \text{var}(\hat{\alpha}) + \frac{(\alpha - \log(-\log(1-0.5)))^2}{\beta^4} \text{var}(\hat{\beta}) + 2\frac{\log(-\log(1-0.5)) - \alpha}{\beta^3} \text{cov}(\hat{\alpha}, \hat{\beta})$$

The asymptotic $(1 - \alpha)100\%$ CI of x_0 is $[\hat{x}_0 - z_{\alpha/2} \sqrt{\text{var}(\hat{x}_0)}, \hat{x}_0 + z_{\alpha/2} \sqrt{\text{var}(\hat{x}_0)}]$

Now, I will calculate these values using the following codes.

```
# LD50 point est, logit
x0 = - glm_logit$coefficients[1]/glm_logit$coefficients[2]
round(exp(x0), 3)
```

```
## (Intercept)
```

```
##      7.389
```

```
# 95% CI
beta0 = glm_logit$coefficients[1]
beta1 = glm_logit$coefficients[2]
betacov = vcov(glm_logit) # inverse fischer information
varx0 = betacov[1, 1]/(beta1^2) + betacov[2, 2]*(beta0^2)/(beta1^4) - 2*betacov[1,2]*beta0/(beta1^3)
se = sqrt(varx0)
round(exp(x0 + c(qnorm(0.05), -qnorm(0.05)) * sqrt(varx0)), 3)
```

```
## [1] 5.51 9.91
```

```

# LD50 point est, probit
x0 = - glm_probit$coefficients[1]/glm_probit$coefficients[2]
round(exp(x0), 3)

## (Intercept)
##          7.436

# 95% CI
beta0 = glm_probit$coefficients[1]
beta1 = glm_probit$coefficients[2]
betacov = vcov(glm_probit) # inverse fischer information
varx0 = betacov[1, 1]/(beta1^2) + betacov[2, 2]*(beta0^2)/(beta1^4) - 2*betacov[1,2]*beta0/(beta1^3)
se = sqrt(varx0)
round(exp(x0 + c(qnorm(0.05), -qnorm(0.05)) * sqrt(varx0)), 3)

## [1] 5.583 9.904

# LD50 point est, cloglog
x0 = (log(-log(1 - 0.5)) - glm_cloglog$coefficients[1])/(glm_cloglog$coefficients[2])
round(exp(x0), 3)

## (Intercept)
##          8.841

# 95% CI
beta0 = glm_cloglog$coefficients[1]
beta1 = glm_cloglog$coefficients[2]
betacov = vcov(glm_cloglog) # inverse fischer information
varx0 = betacov[1, 1]/(beta1^2) + betacov[2, 2]*(beta0 - (log(-log(1 - 0.5))))^2/(beta1^4) + 2*betacov[1,2]*(beta0 - (log(-log(1 - 0.5))))/(beta1^3)
se = sqrt(varx0)
round(exp(x0 + c(qnorm(0.05), -qnorm(0.05)) * sqrt(varx0)), 3)

## [1] 6.526 11.977

```

The results are as follows.

Model	Estimate LD50	90% CI
logit	7.389	(5.510-9.910)
probit	7.436	(5.583-9.904)
c-log-log	8.841	(6.526-11.977)

2

- Amount: one-time two-year scholarship
- Offer: the number of offers made with the corresponding scholarship
- Enrolls: the number of offer accepted

```

# load data
amount = seq(10, 90, 5)
offers = c(4, 6, 10, 12, 39, 36, 22, 14, 10, 12, 8, 9, 3, 1, 5, 2, 1)
enrolls = c(0, 2, 4, 2, 12, 14, 10, 7, 5, 5, 3, 5, 2, 0, 4, 2, 1)

data_2 = data.frame(amount, offers, enrolls)

# visualization
# plot(data_2$amount, data_2$enrolls/data_2$offers, xlab = 'amount', ylab = 'Proportion enrollment', ce

# data prep
x = data_2$amount
y = data_2$enrolls
m = data_2$offers + data_2$enrolls
resp = cbind(y, m-y)

```

(a) How does the model fit the data?

```

# fit logistic model, logit
glm_logit = glm(resp ~ x, family = binomial(link = 'logit'))

```

I employed the logistic regression model to investigate the relationship between the scholarship amount and enrollment rate.

Model: $g(P(enrolls)) = \beta_0 + \beta_1 X$

By the rule of thumb, there are not sufficient people in each scholarship group. Therefore, I will conduct Hosmer-Lemeshow test to check the goodness of fit.

```

# HL test
hltest(glm_logit)

```

```

##
##      The Hosmer-Lemeshow goodness-of-fit test
##
##  Group Size Observed  Expected
##      1   26         6  5.477335
##      2   14         2  3.264660
##      3   51        12 12.651525
##      4   50        14 13.178563
##      5   32        10  8.949818
##      6   21         7  6.223997
##      7   32        10 10.339060
##      8   30        10 10.943589
##      9   16         7  6.971452
##
##          Statistic =  1.26951
## degrees of freedom =  7
##          p-value =  0.98924

```

P-value of 0.990 indicates that the model fits well to the data.

(b) How do you interpret the relationship between the scholarship amount and enrollment rate? What is 95% CI?

```
summary(glm_logit)

##
## Call:
## glm(formula = resp ~ x, family = binomial(link = "logit"))
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.597627   0.365184  -4.375 1.22e-05 ***
## x            0.016290   0.007893   2.064  0.039 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 9.0736  on 16  degrees of freedom
## Residual deviance: 4.8285  on 15  degrees of freedom
## AIC: 52.086
##
## Number of Fisher Scoring iterations: 4

# 95% CI for beta, logit
beta = glm_logit$coefficients[2]
se = sqrt(vcov(glm_logit)[2, 2])
round(beta + c(qnorm(0.025), -qnorm(0.025)) * se, 3)

## [1] 0.001 0.032
```

The model assumes that when the amount scholarship is zero, the log odds for enrollment among offer is -1.598. One unit increase of scholarship amount (in thousand dollars) will increase the log odds of enrollment by 0.0163. With 95% confidence, the true β_1 lies within (0.001-0.032).

(c) How much scholarship should we provide to get 40% yield rate (the percentage of admitted students who enroll?) What is the 95% CI?